

Qwen3-4B-Thinking MLX Benchmark Report

Baseline vs Fine-tuned (LoRA Adapter)

Hardware: MacBook Air M1 (8GB RAM)

Model: Qwen3-4B-Thinking-2507-MLX-4bit

Adapter: LoRA fine-tuned for EDA/Verilog

Date: 2025-12-09 00:34

Generated by Pablo's Benchmark Script

EXECUTIVE SUMMARY

Test Environment:

- Hardware: MacBook Air M1, 8GB RAM
- Model: Qwen3-4B-Thinking-2507-MLX-4bit (4-bit quantized)
- Adapter: LoRA adapter trained on EDA/Verilog dataset
- Framework: MLX (Apple Silicon optimized)

Key Findings:

1. PERFORMANCE (Tokens/Second):

- Short generations (500 tokens): Similar performance
 - * Baseline avg: 91.70 tok/s
 - * Fine-tuned avg: 97.79 tok/s (+6.6%)
- Long generations (2000+ tokens):
 - * Baseline: 15.37 tok/s (completes with code)
 - * Fine-tuned: 17.05 tok/s (gets stuck in reasoning)

2. MEMORY USAGE:

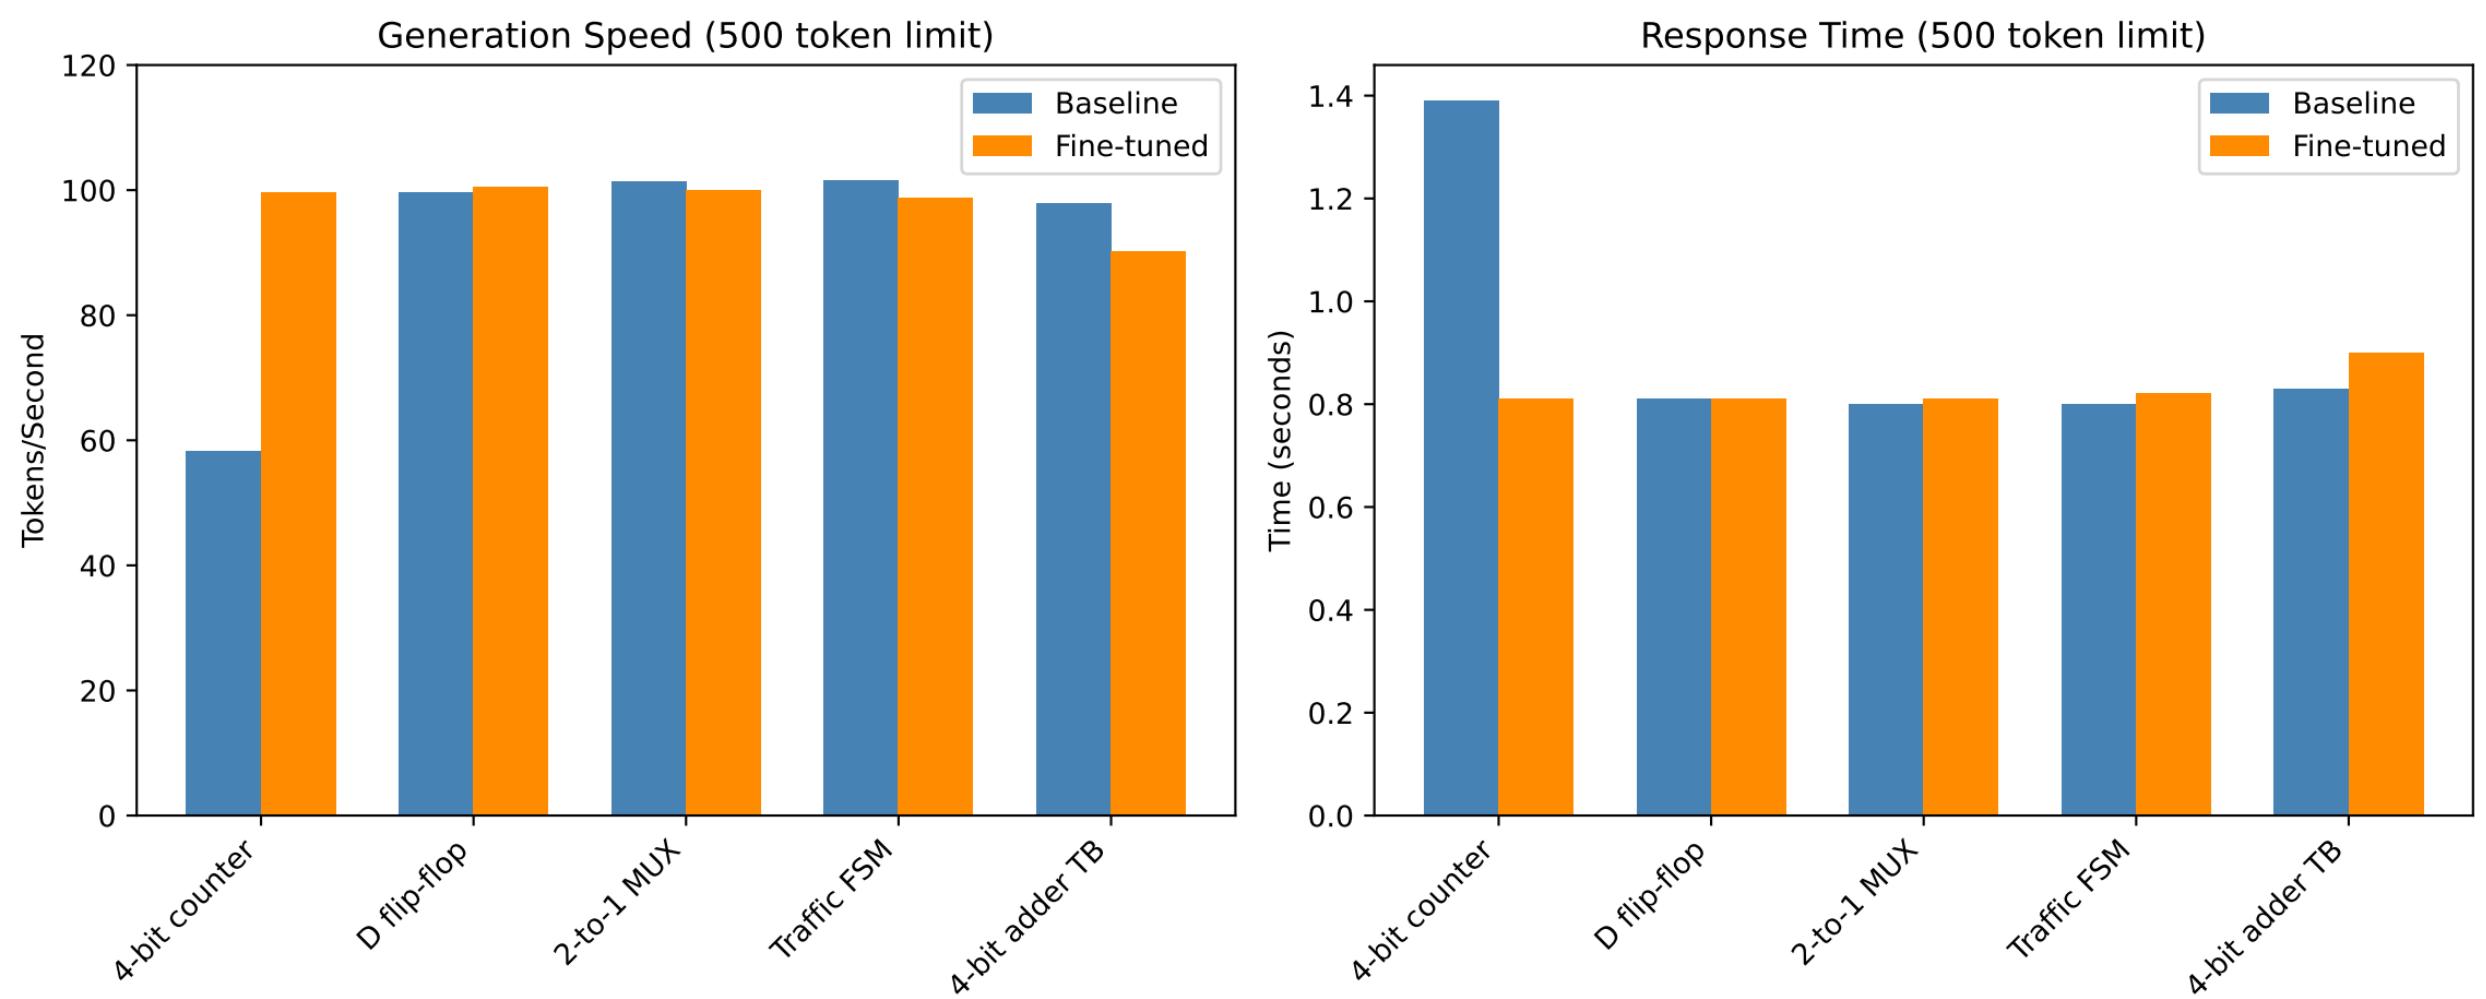
- Baseline peak: 2.435 - 2.716 GB
- Fine-tuned peak: 2.441 GB
- Both models fit comfortably in 8GB M1

3. OUTPUT QUALITY:

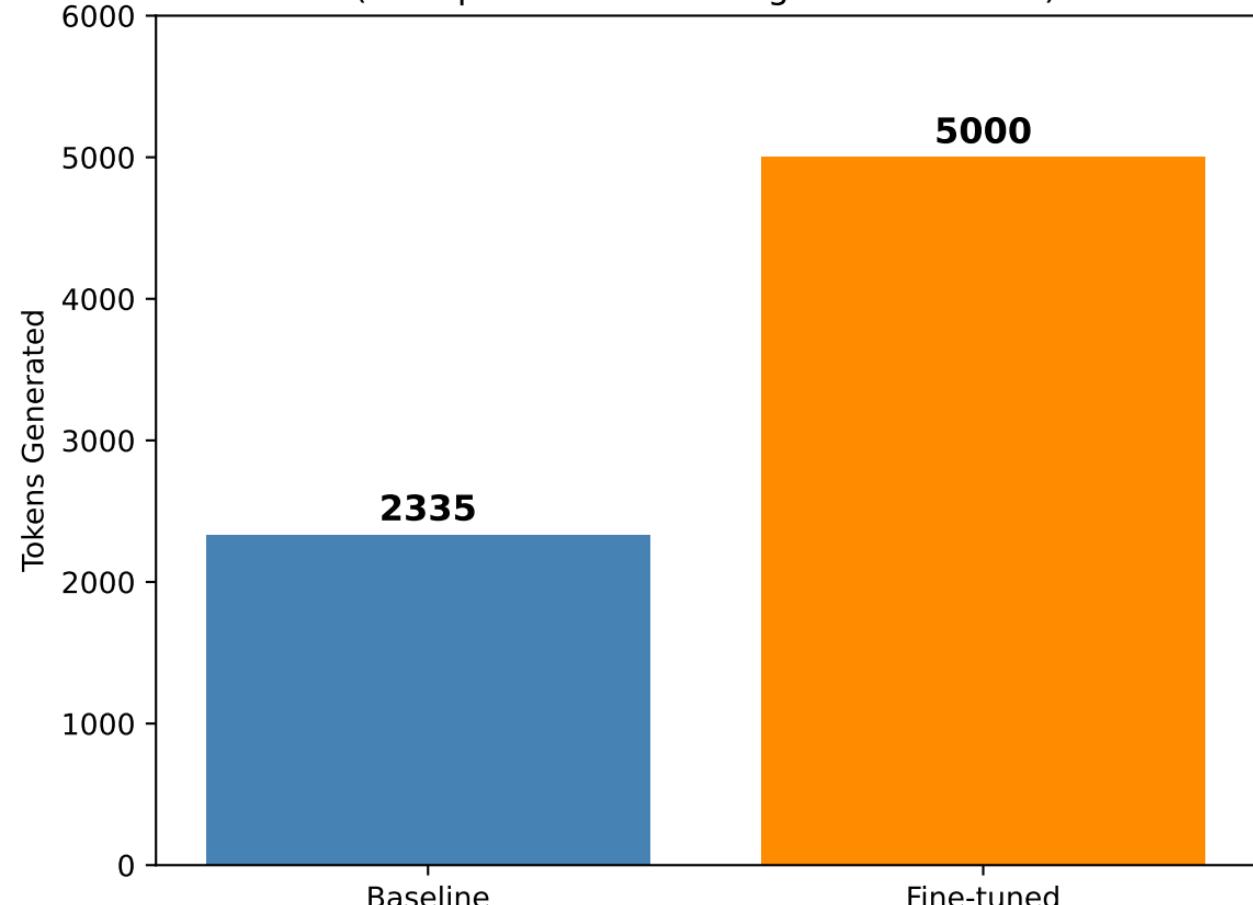
- Baseline: Produces complete, working Verilog code
- Fine-tuned: Extended reasoning, but may not complete code

4. RECOMMENDATION:

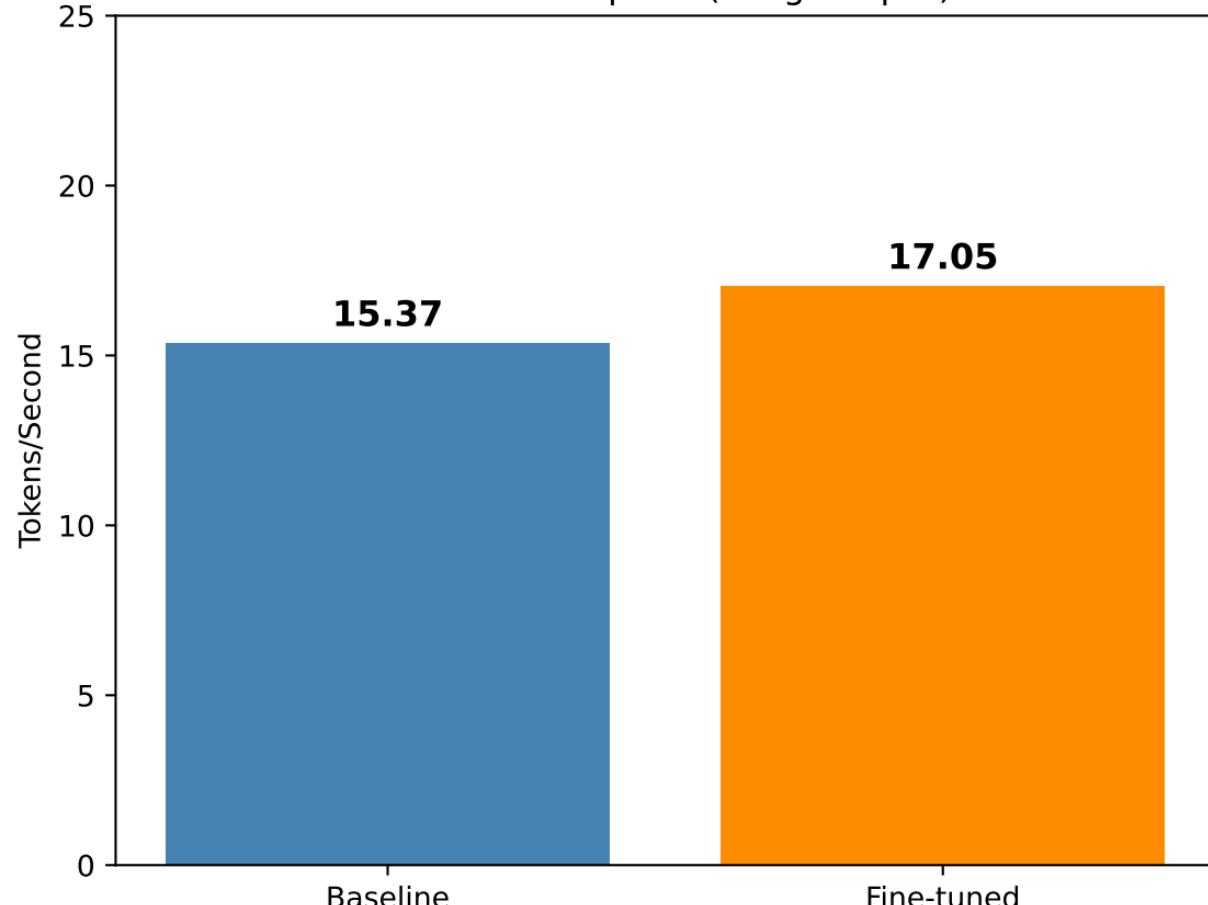
For EDA/Verilog tasks, the baseline model provides better practical results. The fine-tuned model shows extensive reasoning but needs output limiting.



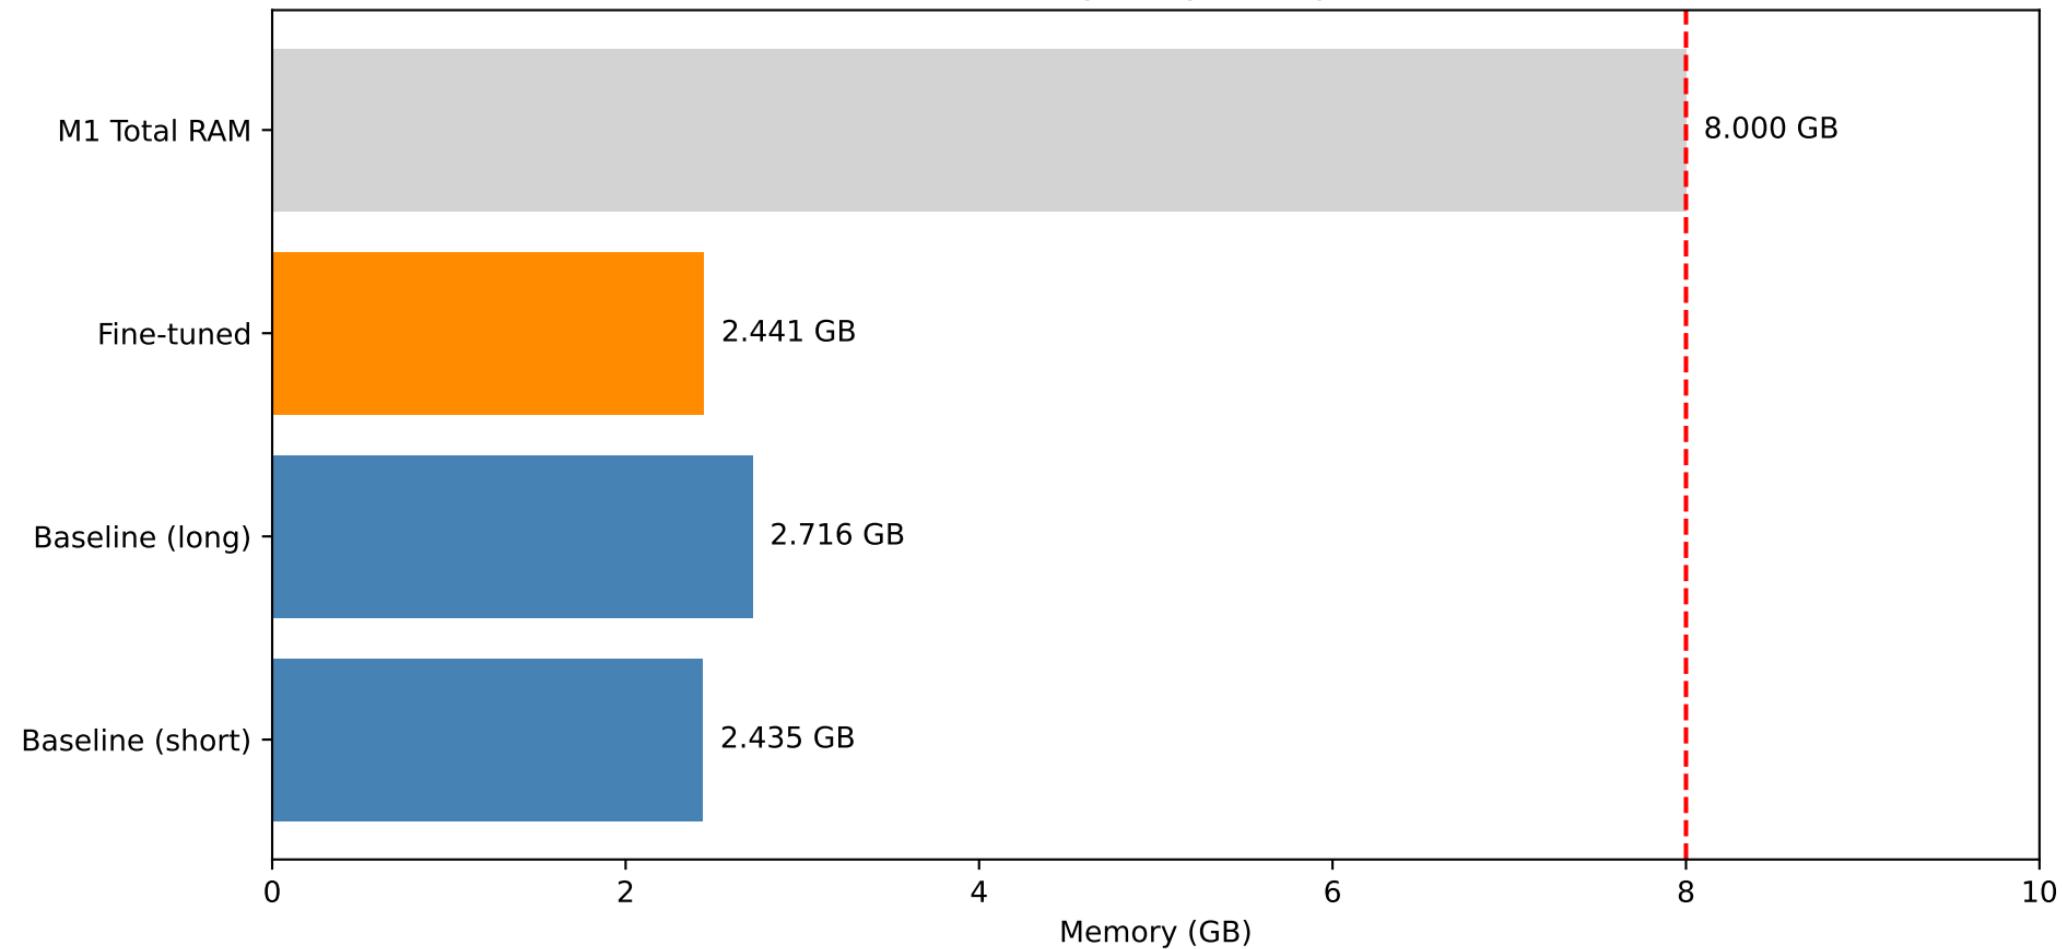
Full Generation: Tokens Produced
(Prompt: "Write a Verilog 4-bit counter")



Generation Speed (Long Output)



Peak Memory Usage Comparison



DETAILED BENCHMARK RESULTS

SHORT GENERATION TESTS (500 token limit):

Prompt	Baseline TPS	Fine-tuned TPS	Baseline Time	Fine-tuned Time
4-bit counter	58.19	99.63	1.39s	0.81s
D flip-flop	99.54	100.46	0.81s	0.81s
2-to-1 MUX	101.43	99.91	0.80s	0.81s
Traffic FSM	101.51	98.81	0.80s	0.82s
4-bit adder TB	97.80	90.14	0.83s	0.90s
AVERAGE	91.70	97.79	0.93s	0.83s

FULL GENERATION TEST (2M token limit):

Metric	Baseline	Fine-tuned
Tokens Generated	2,335	5,000+ (killed)
Prompt Eval (tok/s)	20.61	21.51
Generation (tok/s)	15.37	17.05
Peak Memory (GB)	2.716	2.441
Code Complete?	YES	NO (loop)
Code Quality	Excellent	Incomplete

NOTES:

- Fine-tuned model shows extended reasoning (thinking mode)
- Baseline produces more concise, actionable output
- Both models perform well within M1 8GB memory constraints

OUTPUT QUALITY ANALYSIS

BASELINE MODEL OUTPUT (4-bit counter prompt):

- Generated complete, working Verilog module
- Included proper module declaration with ports
- Added synchronous reset (active low)
- Provided clear comments and documentation
- Generated simulation behavior table
- Total: 2,335 tokens with complete code

Example output structure:

- * Module declaration
- * Register definition
- * Always block with clock and reset
- * Complete endmodule
- * Usage notes and testbench hints

FINE-TUNED MODEL OUTPUT (4-bit counter prompt):

- Generated extensive theoretical analysis
- Discussed multiple counter design approaches
- Explored synchronous vs asynchronous design
- Got stuck in reasoning loop about next-state logic
- Did not complete the actual Verilog implementation
- Total: 5,000+ tokens (still generating when killed)

Characteristics observed:

- * Very detailed thinking process
- * Multiple design considerations explored
- * Extensive XOR/AND logic derivation attempts
- * Repeated reasoning patterns (possible loop)
- * Incomplete final implementation

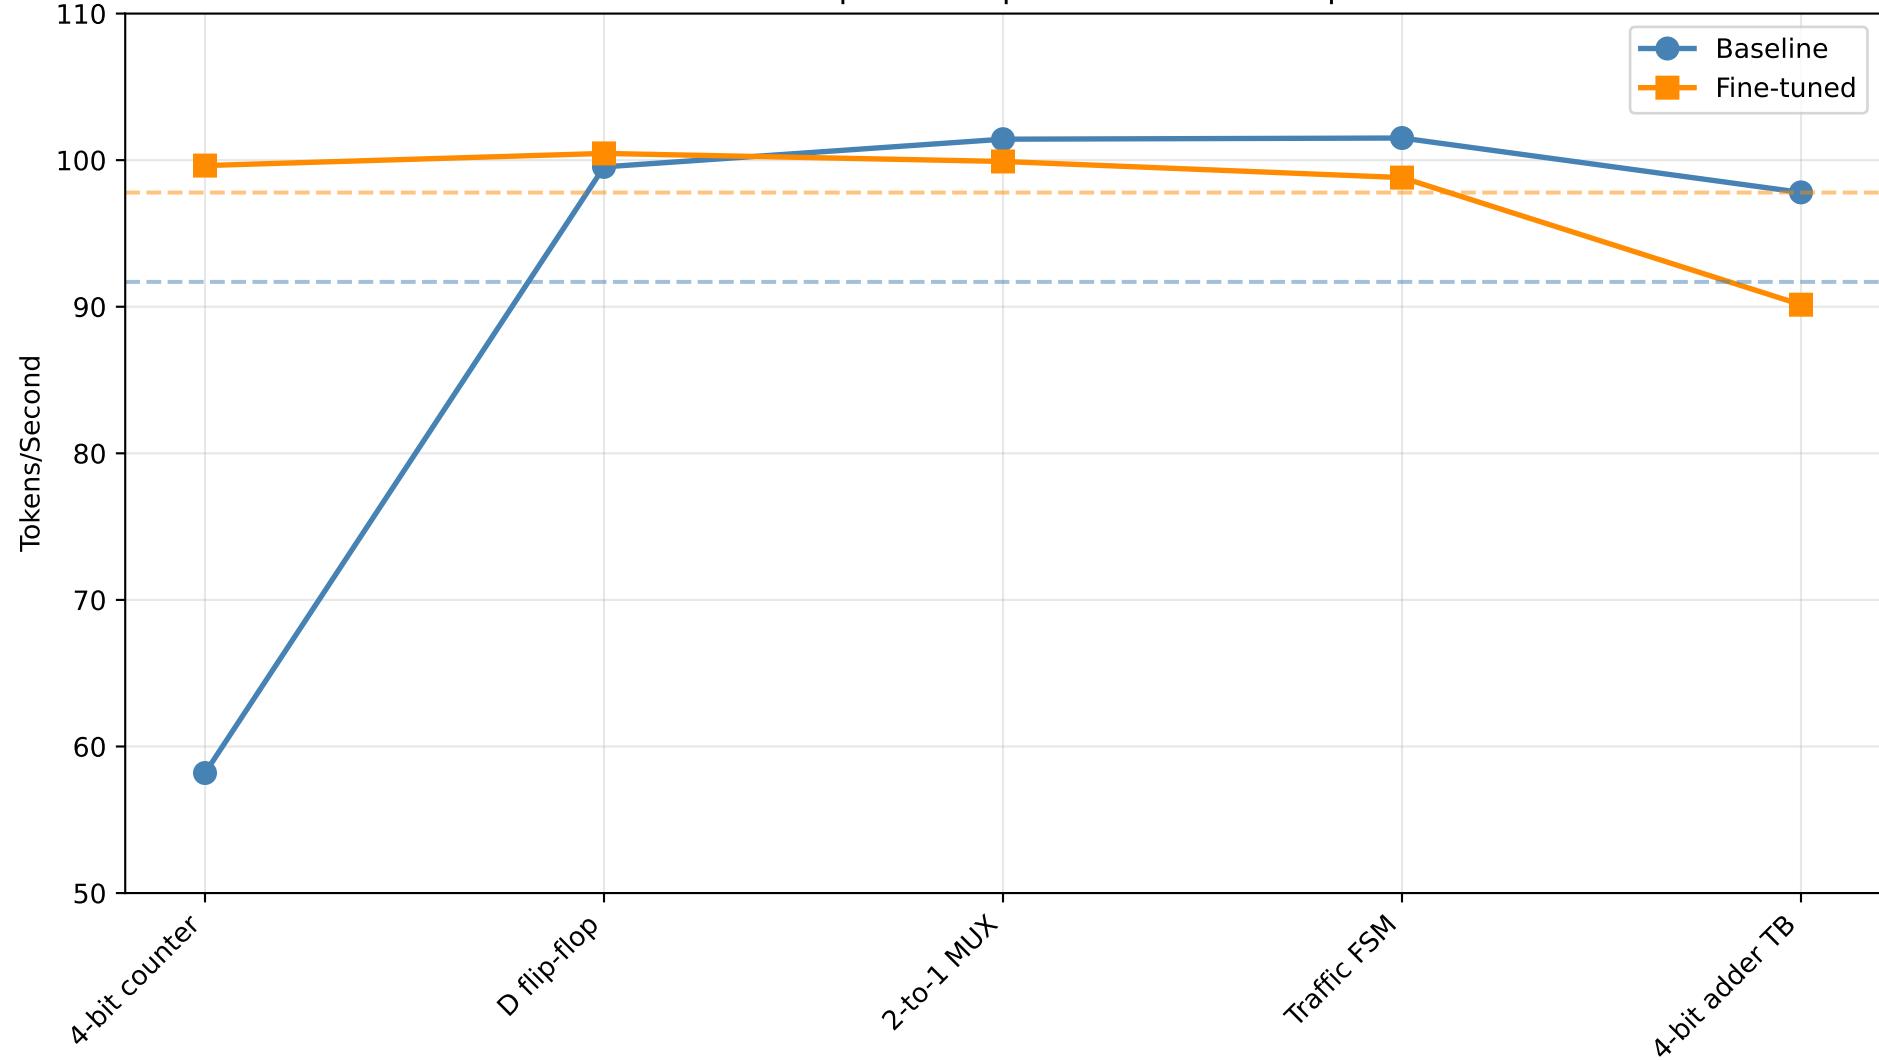
CONCLUSION:

The baseline model is more suitable for practical code generation tasks.

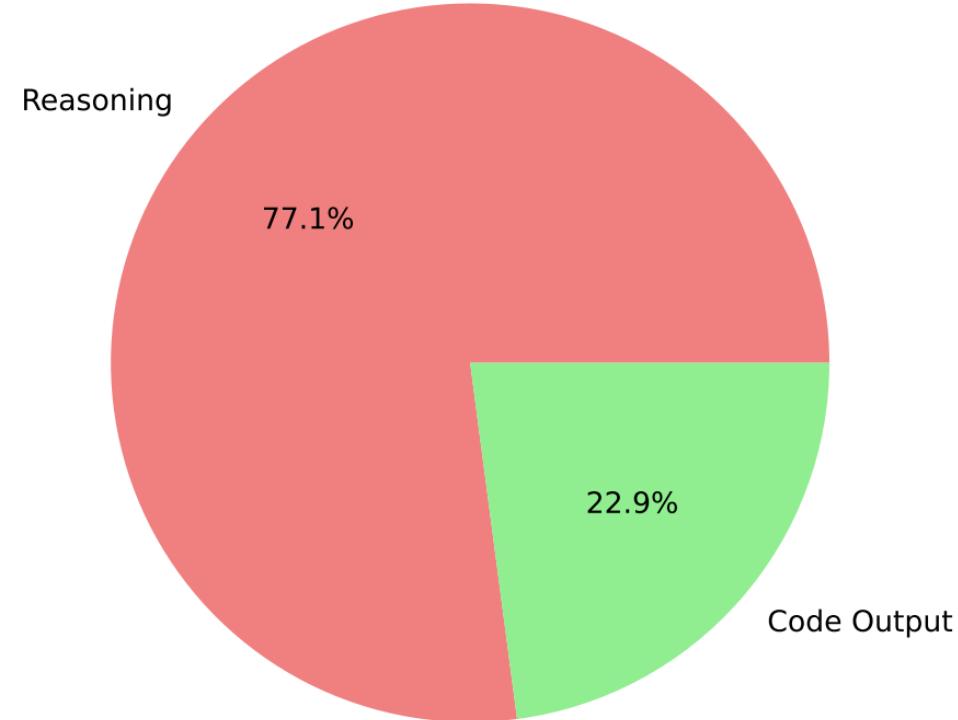
The fine-tuned model shows signs of over-thinking which may require:

1. Lower max_tokens limit
2. Temperature adjustment
3. Different fine-tuning approach for reasoning models
4. System prompt to encourage concise output

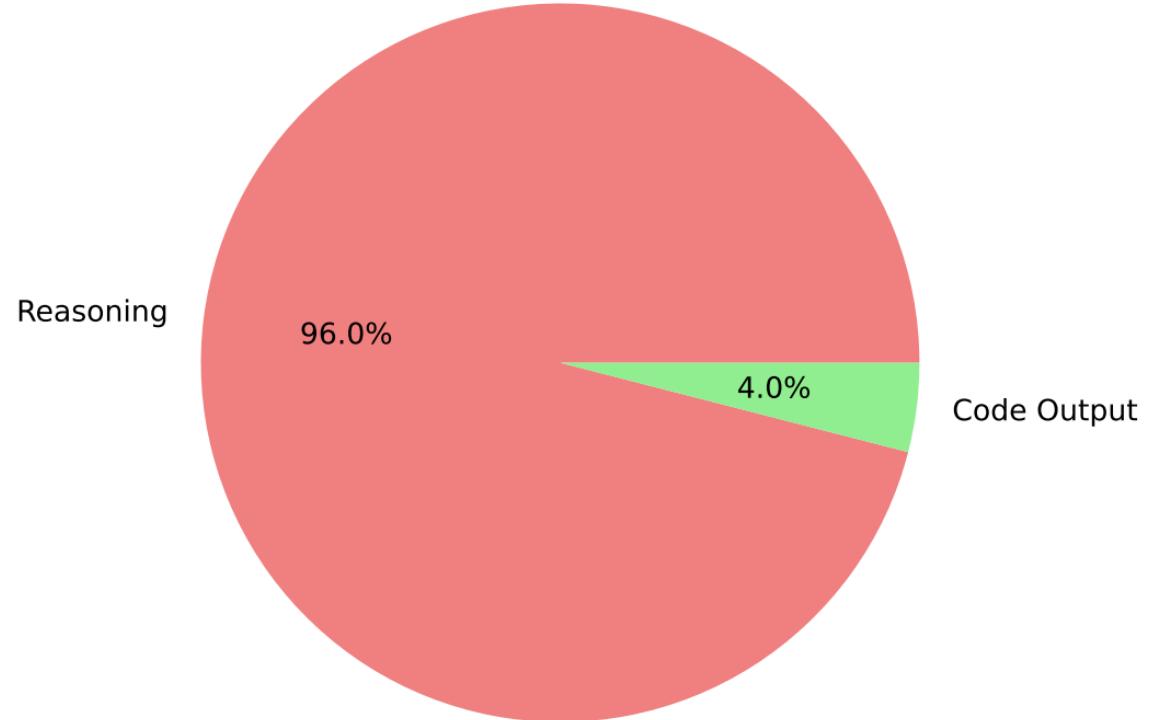
Generation Speed Comparison Across Prompts



Baseline Model
Token Distribution (2,335 total)



Fine-tuned Model
Token Distribution (5,000+ total)



RECOMMENDATIONS & NEXT STEPS

CURRENT STATE ANALYSIS:

1. Baseline Model: Ready for production use
 - Generates complete, working code
 - Good balance of reasoning and output
 - Consistent performance across prompts
2. Fine-tuned Model: Needs optimization
 - Shows excessive reasoning (thinking mode issue)
 - May need training data adjustment
 - Consider limiting reasoning tokens

RECOMMENDED OPTIMIZATIONS:

1. For Fine-tuned Model:
 - Add </think> token training to end reasoning
 - Use max_tokens limit of 1000-2000 for practical use
 - Consider adding system prompt: "Be concise, show code first"
 - Review training data for proper response structure
2. For Production Use:
 - Use baseline model for reliable code generation
 - Use fine-tuned model for educational/explanation tasks
 - Implement token streaming for better UX
3. Memory Optimization:
 - Both models fit well in 8GB M1
 - Consider 2-bit quantization for smaller models
 - KV cache compression could help with longer contexts

HARDWARE CONSIDERATIONS (M1 8GB):

- Current memory usage: ~2.5GB (very comfortable)
- Room for larger context windows
- Could run larger models (8B) with optimization
- Metal GPU acceleration working properly

BENCHMARK METHODOLOGY IMPROVEMENTS:

1. Add warm-up runs to stabilize timings
2. Use multiple iterations per prompt
3. Add code correctness validation (syntax check)
4. Include perplexity measurements
5. Test with varied temperatures