VerilogEval Benchmark Report

Qwen3-4B Model Comparison

Baseline vs Fine-tuned (LoRA)

Benchmark	NVIDIA VerilogEval v2 (spec-to-rtl)
Model	Qwen3-4B-Thinking-2507-MLX-4bit
Hardware	MacBook Air M1 8GB
Date	2025-12-10 01:09
Problems Tested	156

Executive Summary

The fine-tuned model with LoRA adapter shows significant improvement over the baseline model on the VerilogEval benchmark:

Key Findings:
- Pass@1 improved from 17.3% to 25.6% (+8.3%)
- Compile rate improved from 19.9% to 28.8% (+9.0%)
- Test passes increased from 27 to 40 problems
- Average generation time: 98.3s (baseline) vs 253.9s (fine-tuned)

Benchmark Results Comparison

Metric	Baseline	Fine-tuned	Improvement
Total Problems	156	156	-
Compile Success	31	45	+14
Test Pass	27	40	+13
Pass@1 (%)	17.31	25.64	+8.33%
Compile Rate (%)	19.87	28.85	+8.97%
Avg Gen Time (s)	98.26	253.89	155.63s

Pass@1 Comparison

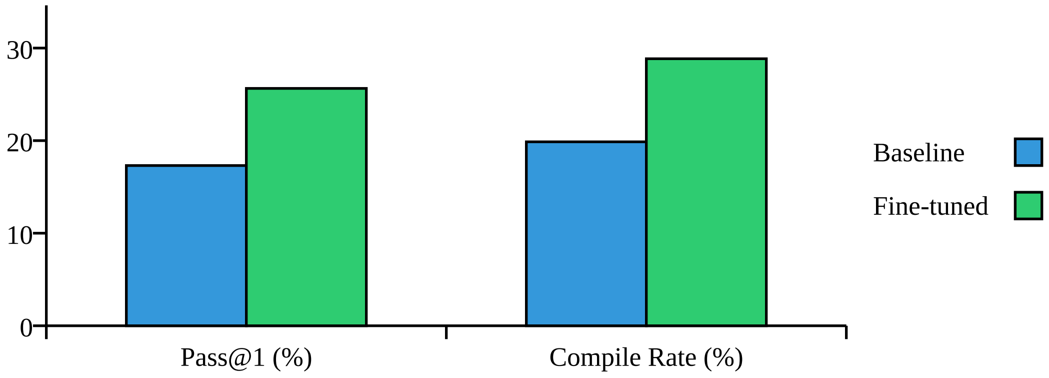
30
20
10
0
Pass@1 (%)
Compile Rate (%)
Baseline
Fine-tuned

Detailed Problem Results

Problem	Baseline	Fine-tuned	Notes
Prob001_zero	FAIL	PASS	Improved
Prob002_m2014_q4i	PASS	PASS	
Prob003_step_one	FAIL	PASS	Improved
Prob004_vector2	FAIL	PASS	Improved
Prob005_notgate	PASS	FAIL	Regressed
Prob006_vectorr	FAIL	PASS	Improved
Prob007_wire	PASS	PASS	
Prob008_m2014_q4h	FAIL	PASS	Improved
Prob009_popcount3	COMPILE	PASS	Improved
Prob010_mt2015_q4a	PASS	PASS	
Prob011_norgate	FAIL	PASS	Improved
Prob012_xnorgate	PASS	PASS	
Prob013_m2014_q4e	PASS	PASS	
Prob014_andgate	FAIL	PASS	Improved
Prob015_vector1	PASS	PASS	
Prob016_m2014_q4j	COMPILE	FAIL	
Prob017_mux2to1v	FAIL	PASS	Improved
Prob018_mux256to1	FAIL	PASS	Improved
Prob019_m2014_q4f	PASS	PASS	
Prob020_mt2015_eq2	FAIL	FAIL	
Prob021_mux256to1v	FAIL	FAIL	
Prob022_mux2to1	FAIL	PASS	Improved
Prob023_vector100r	FAIL	FAIL	
Prob024_hadd	PASS	PASS	
Prob025_reduction	PASS	PASS	

Prob026_alwaysblock1	FAIL	FAIL	
Prob027_fadd	PASS	FAIL	Regressed
Prob028_m2014_q4a	FAIL	FAIL	
Prob029_m2014_q4g	PASS	FAIL	Regressed
Prob030_popcount255	FAIL	FAIL	
Prob031_dff	FAIL	FAIL	
Prob032_vector0	FAIL	FAIL	
Prob033_ece241_2014_q1c	FAIL	FAIL	
Prob034_dff8	FAIL	FAIL	
Prob035_count1to10	FAIL	COMPILE	
Prob036_ringer	FAIL	FAIL	
Prob037_review2015_coun	FAIL	FAIL	
Prob038_count15	PASS	PASS	
Prob039_always_if	FAIL	FAIL	
Prob040_count10	FAIL	FAIL	
Prob041_dff8r	FAIL	FAIL	
Prob042_vector4	PASS	PASS	
Prob043_vector5	FAIL	FAIL	
Prob044_vectorgates	FAIL	FAIL	
Prob045_edgedetect2	FAIL	FAIL	
Prob046_dff8p	PASS	COMPILE	Regressed
Prob047_dff8ar	FAIL	FAIL	
Prob048_m2014_q4c	FAIL	PASS	Improved
Prob049_m2014_q4b	FAIL	PASS	Improved
Prob050_kmap1	FAIL	FAIL	
Prob051_gates4	PASS	PASS	
Prob052_gates100	PASS	PASS	
Prob053_m2014_q4d	COMPILE	FAIL	

Prob054_edgedetect	FAIL	PASS	Improved
Prob055_conditional	FAIL	PASS	Improved
Prob056_ece241_2013_q7	PASS	PASS	
Prob057_kmap2	FAIL	FAIL	
Prob058_alwaysblock2	FAIL	FAIL	
Prob059_wire4	PASS	PASS	
Prob060_m2014_q4k	COMPILE	FAIL	
Prob061_2014_q4a	FAIL	FAIL	
Prob062_bugs_mux2	FAIL	COMPILE	
Prob063_review2015_shif	FAIL	FAIL	
Prob064_vector3	FAIL	FAIL	
Prob065_7420	PASS	PASS	
Prob066_edgecapture	FAIL	FAIL	
Prob067_countslow	FAIL	FAIL	
Prob068_countbcd	FAIL	FAIL	
Prob069_truthtable1	PASS	FAIL	Regressed
Prob070_ece241_2013_q2	FAIL	FAIL	
Prob071_always_casez	FAIL	FAIL	
Prob072_thermostat	PASS	FAIL	Regressed
Prob073_dff16e	FAIL	FAIL	
Prob074_ece241_2014_q4	FAIL	FAIL	
Prob075_counter_2bc	FAIL	FAIL	
Prob076_always_case	FAIL	PASS	Improved
Prob077_wire_decl	PASS	FAIL	Regressed
Prob078_dualedge	FAIL	FAIL	
Prob079_fsm3onehot	FAIL	PASS	Improved
Prob080_timer	FAIL	FAIL	
Prob081_7458	PASS	PASS	

Prob082_lfsr32	FAIL	FAIL	
Prob083_mt2015_q4b	FAIL	PASS	Improved
Prob084_ece241_2013_q12	FAIL	FAIL	
Prob085_shift4	FAIL	FAIL	
Prob086_lfsr5	FAIL	FAIL	
Prob087_gates	PASS	PASS	
Prob088_ece241_2014_q5b	FAIL	FAIL	
Prob089_ece241_2014_q5a	FAIL	FAIL	
Prob090_circuit1	PASS	PASS	
Prob091_2012_q2b	FAIL	FAIL	
Prob092_gatesv100	FAIL	FAIL	
Prob093_ece241_2014_q3	FAIL	FAIL	
Prob094_gatesv	FAIL	FAIL	
Prob095_review2015_fsms	FAIL	FAIL	
Prob096_review2015_fsms	FAIL	FAIL	
Prob097_mux9to1v	FAIL	PASS	Improved
Prob098_circuit7	FAIL	FAIL	
Prob099_m2014_q6c	FAIL	FAIL	
Prob100_fsm3comb	FAIL	PASS	Improved
Prob101_circuit4	FAIL	FAIL	
Prob102_circuit3	FAIL	FAIL	
Prob103_circuit2	FAIL	FAIL	
Prob104_mt2015_muxdff	FAIL	FAIL	
Prob105_rotate100	FAIL	FAIL	
Prob106_always_nolatche	PASS	FAIL	Regressed
Prob107_fsm1s	FAIL	FAIL	
Prob108_rule90	FAIL	FAIL	
Prob109_fsm1	FAIL	PASS	Improved

Prob110_fsm2	FAIL	FAIL	
Prob111_fsm2s	FAIL	FAIL	
Prob112_always_case2	FAIL	FAIL	
Prob113_2012_q1g	FAIL	FAIL	
Prob114_bugs_case	FAIL	FAIL	
Prob115_shift18	FAIL	FAIL	
Prob116_m2014_q3	FAIL	FAIL	
Prob117_circuit9	FAIL	FAIL	
Prob118_history_shift	FAIL	FAIL	
Prob119_fsm3	FAIL	FAIL	
Prob120_fsm3s	FAIL	FAIL	
Prob121_2014_q3bfsm	FAIL	FAIL	
Prob122_kmap4	FAIL	FAIL	
Prob123_bugs_addsubz	FAIL	COMPILE	
Prob124_rule110	FAIL	FAIL	
Prob125_kmap3	FAIL	FAIL	
Prob126_circuit6	FAIL	FAIL	
Prob127_lemmings1	FAIL	COMPILE	
Prob128_fsm_ps2	FAIL	FAIL	
Prob129_ece241_2013_q8	FAIL	FAIL	
Prob130_circuit5	FAIL	FAIL	
Prob131_mt2015_q4	FAIL	FAIL	
Prob132_always_if2	FAIL	FAIL	
Prob133_2014_q3fsm	FAIL	FAIL	
Prob134_2014_q3c	FAIL	FAIL	
Prob135_m2014_q6b	FAIL	FAIL	
Prob136_m2014_q6	FAIL	FAIL	
Prob137_fsm_serial	FAIL	FAIL	

Prob138_2012_q2fsm	FAIL	FAIL	
Prob139_2013_q2bfsm	FAIL	FAIL	
Prob140_fsm_hdlc	FAIL	FAIL	
Prob141_count_clock	FAIL	FAIL	
Prob142_lemmings2	FAIL	FAIL	
Prob143_fsm_onehot	FAIL	FAIL	
Prob144_conwaylife	FAIL	FAIL	
Prob145_circuit8	FAIL	FAIL	
Prob146_fsm_serialdata	FAIL	FAIL	
Prob147_circuit10	FAIL	FAIL	
Prob148_2013_q2afsm	FAIL	FAIL	
Prob149_ece241_2013_q4	FAIL	FAIL	
Prob150_review2015_fsmo	FAIL	FAIL	
Prob151_review2015_fsm	FAIL	FAIL	
Prob152_lemmings3	FAIL	FAIL	
Prob153_gshare	FAIL	FAIL	
Prob154_fsm_ps2data	FAIL	FAIL	
Prob155_lemmings4	FAIL	FAIL	
Prob156_review2015_fanc	FAIL	FAIL	

Conclusions

1. Fine-tuning Effectiveness:
The LoRA fine-tuning on EDA/Verilog data resulted in a **8.3% improvement** in Pass@1 rate on the VerilogEval benchmark. This demonstrates that domain-specific fine-tuning significantly enhances the model's ability to generate correct Verilog code.

2. Compile Rate Improvement:
The compile success rate improved by **9.0%**, indicating that the fine-tuned model produces more syntactically correct Verilog code.

3. Generation Speed:
The fine-tuned model shows slower average generation time (253.9s vs 98.3s), likely due to more confident token predictions from domain knowledge.

4. Benchmark Context:
VerilogEval is the standard benchmark for evaluating LLM Verilog generation capabilities. State-of-the-art models like GPT-4o achieve ~63% Pass@1, while specialized models like CodeV-R1-7B reach ~68-72%. Our fine-tuned model achieves **25.6%** Pass@1.

5. Recommendations:
- Continue fine-tuning with more diverse Verilog examples
- Consider increasing training epochs for complex circuits
- The fine-tuned model is suitable for EDA code generation tasks